

Bayesian networks as networks of mechanisms

Simon Frost

Table of contents

1	Overview	1
2	Setup	2
3	Variables, states and mechanisms	2
4	Validation	5
5	The derived graph	5
6	Why “not just a DAG”	7
7	Binding tables	8
8	Brute-force evaluation	8
9	Summary	10
10	References	10

1 Overview

A Bayesian network is usually introduced as a directed acyclic graph whose nodes carry conditional probability tables ([Koller and Friedman 2009](#)). `BayesianNetworks.jl` takes a slightly different view, the one that makes networks composable: a network is a collection of *variables*, each with an ordered list of *states*, and a collection of *mechanisms*, each of which generates one variable from an ordered list of input variables. The graph is derived from the mechanisms; it is not stored. Treating the mechanisms rather than the arrows as the data is what lets a network with some variables left ungenerated make sense, which is the categorical reading of a Bayesian network as a causal theory ([Fong 2012](#)).

This vignette builds a reference ecological network – climate and irrigation drive soil moisture, soil moisture and grazing pressure drive vegetation, and vegetation determines habitat quality

and hence occupancy (the running example of the package specification) – then inspects it, validates it and draws it. The companion vignettes cover the wiring-diagram view and evaluation, observation versus intervention, open networks and composition, serialisation, dynamic networks and model cards.

2 Setup

The package re-exports what it needs from `ACSets` and `FiniteKernels.jl`, so this one using is enough.

```
using BayesianNetworks
```

3 Variables, states and mechanisms

The `bayesnet` function takes `name ⇒ states` pairs and a list of `target ⇒ parents` mechanisms. Parent order matters: it fixes the axis order of the conditional probability table that will later be bound to the mechanism.

```
bn = bayesnet(:Climate ⇒ [:dry, :normal, :wet],
              :Irrigation ⇒ [:low, :high],
              :SoilMoisture ⇒ [:low, :medium, :high],
              :GrazingPressure ⇒ [:low, :high],
              :Vegetation ⇒ [:sparse, :moderate, :dense],
              :HabitatQuality ⇒ [:poor, :good],
              :Occupancy ⇒ [:absent, :present];
              mechanisms = [:SoilMoisture ⇒ (:Climate, :Irrigation),
                            :Vegetation ⇒ (:SoilMoisture, :GrazingPressure),
                            :HabitatQuality ⇒ :Vegetation,
                            :Occupancy ⇒ :HabitatQuality])
```

Variable	variable_name	space_ref
1	Climate	NoRef()
2	Irrigation	NoRef()
3	SoilMoisture	NoRef()
4	GrazingPressure	NoRef()
5	Vegetation	NoRef()
6	HabitatQuality	NoRef()
7	Occupancy	NoRef()

State	state_variable	state_name	state_position
1	1	dry	1
2	1	normal	2
3	1	wet	3
4	2	low	1
5	2	high	2
6	3	low	1
7	3	medium	2
8	3	high	3
9	4	low	1
10	4	high	2
11	5	sparse	1
12	5	moderate	2
13	5	dense	3
14	6	poor	1
15	6	good	2
16	7	absent	1
17	7	present	2

Mechanism	target	mechanism_name	kernel_ref
1	3	SoilMoisture_mechanism	NoRef()
2	5	Vegetation_mechanism	NoRef()
3	6	HabitatQuality_mechanism	NoRef()
4	7	Occupancy_mechanism	NoRef()
5	1	Climate_mechanism	NoRef()
6	2	Irrigation_mechanism	NoRef()
7	4	GrazingPressure_mechanism	NoRef()

Input	input_mechanism	input_variable	input_position
1	1	1	1
2	1	2	2
3	2	3	1
4	2	4	2
5	3	5	1
6	4	6	1

The same network is available as `reference_habitat_bn()`. Every variable that is not named as a target gets a mechanism without inputs (a prior), so the network is *closed*: each variable has exactly one generating mechanism.

```
mechanism_names(bn)
```

```
7-element Vector{Symbol}:
 :SoilMoisture_mechanism
 :Vegetation_mechanism
 :HabitatQuality_mechanism
 :Occupancy_mechanism
 :Climate_mechanism
 :Irrigation_mechanism
 :GrazingPressure_mechanism
```

The network is an attributed C-set (ACSet) on the schema SchBayesNet ([Patterson et al. 2022](#)), with four objects. There is no Edge object: an arrow from Climate to SoilMoisture is an Input part saying that Climate is the first argument of the mechanism generating SoilMoisture.

```
(variables = nparts(bn, :Variable), states = nparts(bn, :State),
 mechanisms = nparts(bn, :Mechanism), inputs = nparts(bn, :Input))
```

```
(variables = 7, states = 17, mechanisms = 7, inputs = 6)
```

Accessors return part ids; names are recovered with `variable_name`. States and inputs are always returned in their recorded order (`state_position`, `input_position`), never in part-id order.

```
states(bn, :SoilMoisture)
```

```
3-element Vector{Symbol}:
 :low
 :medium
 :high
```

`parents` gives the inputs of the mechanism generating a variable, in the order that fixes the table's axes:

```
variable_name.(Ref(bn), parents(bn, :Vegetation))
```

```
2-element Vector{Symbol}:
 :SoilMoisture
 :GrazingPressure
```

`children` reads the same Input parts the other way, listing every variable whose mechanism consumes the given one:

```
variable_name.(Ref(bn), children(bn, :SoilMoisture))
```

```
1-element Vector{Symbol}:
 :Vegetation
```

4 Validation

`validate` checks the structural rules of the specification: every reference points at an existing part, state and input positions are permutations, no mechanism reads its own target, the derived graph is acyclic, and (with `closed = true`) every variable has exactly one mechanism. It returns nothing when everything is fine and throws a typed exception otherwise; `validation_errors` collects all problems.

```
validate(bn; closed = true, unique_names = true)
```

Let us break the network on a copy: add a second mechanism for Occupancy and an input that closes a cycle.

```
broken = deepcopy(bn)
add_mechanism!(broken, :Occupancy; inputs = [:Climate], name = :second)
add_input!(broken, :Climate_mechanism, :Occupancy)
validation_errors(broken; closed = true)
```

```
2-element Vector{Exception}:
 DuplicateGeneratorError(:Occupancy, 7, [4, 8])
 CyclicBayesNetError([:Climate, :SoilMoisture, :Vegetation, :HabitatQuality,
 :Occupancy], [1, 3, 5, 6, 7])
```

The mutating builders (`add_variable!`, `add_state!`, `add_mechanism!`, `add_input!`) never validate, so intermediate states can be inconsistent; everything else in the package treats networks as immutable values and returns new ones.

5 The derived graph


The DAG is a view: `variable_graph` builds a `Graphs.jl SimpleDiGraph` with one vertex per variable and one edge per input, and `topological_order` sorts the variables so that parents come first.

```
variable_name.(Ref(bn), topological_order(bn))
```

```
7-element Vector{Symbol}:
 :Climate
 :Irrigation
 :SoilMoisture
 :GrazingPressure
 :Vegetation
 :HabitatQuality
 :Occupancy
```


`to_graphviz` draws the same graph. Nodes are labelled with the variable and its states; the result is a `Graphviz.Graph`, the package's own DOT syntax tree, which renders to SVG through the Graphviz binaries.

```
to_graphviz(bn)
```



The layout direction and Graphviz attributes can be changed, and edges can be labelled with the input position, which shows the parent order of `SoilMoisture` and `Vegetation`:

```
to_graphviz(bn; rankdir = "LR", states = false, edge_labels = true)
```



6 Why “not just a DAG”

Three things distinguish this representation from a graph with tables attached.

1. Mechanisms are first-class parts. They have names and kernel references, so an intervention can replace one (the companion vignette on interventions), a composition can identify one across networks, and a refinement can substitute an open sub-network for one (the vignette on open networks). The graph never has to be edited; it follows the mechanisms.
2. Order is explicit. State order and parent order are attributes, so they survive colimits, which renumber parts, and they make the axis convention of the tables unambiguous: a table bound to `Vegetation` has axes (`SoilMoisture`, `GrazingPressure`, `Vegetation`) because that is the recorded input order.
3. Syntax and semantics are separate. The `ACSet` holds no numbers. Kernels live in a `BayesModel` wrapper, keyed by the `kernel_ref` attribute of each mechanism, so the same structure can carry different tables, a different semantics altogether, or none at all.

Canonical forms make the representation comparable: `canonicalize` rennumbers parts deterministically, so two networks built in different orders compare equal.

```

reversed = bayesnet(:Occupancy => [:absent, :present],
  :HabitatQuality => [:poor, :good],
  :Vegetation => [:sparse, :moderate, :dense],
  :GrazingPressure => [:low, :high],
  :SoilMoisture => [:low, :medium, :high],
  :Irrigation => [:low, :high],
  :Climate => [:dry, :normal, :wet];
mechanisms = [:Occupancy => :HabitatQuality,
  :HabitatQuality => :Vegetation,
  :Vegetation => (:SoilMoisture, :GrazingPressure),
  :SoilMoisture => (:Climate, :Irrigation)]
reversed == bn, canonicalize(reversed) == canonicalize(bn)

```

```
(false, true)
```

7 Binding tables

A `BayesModel` wraps the structure with spaces and kernels from `FiniteKernels.jl`. `bind_cpt` takes a table in the parents-first, child-last layout and checks its shape and normalisation against the syntax.

```
m = BayesModel(bn)
m = bind_cpt(m, :Climate ⇒ [0.3, 0.5, 0.2])
m = bind_cpt(m, :HabitatQuality ⇒ [0.85 0.15; 0.4 0.6; 0.15 0.85]) # rows =
Vegetation
missing_kernels(m)
```

```
5-element Vector{Symbol}:
 :SoilMoisture
 :Vegetation
 :Occupancy
 :Irrigation
 :GrazingPressure
```

`reference_habitat_model()` binds all seven tables.

```
m = reference_habitat_model()
validate(m; closed = true, semantics = true)
kernel(m, :HabitatQuality)
```

```
FiniteKernel{Float64}(Vegetation{sparse,moderate,dense} → HabitatQuality{poor,good})
      sparse moderate dense
poor    0.85      0.4    0.15
good    0.15      0.6    0.85
```

8 Brute-force evaluation

With every mechanism bound, the model has a joint distribution: the chain-rule product of the conditional probability tables (Koller and Friedman 2009). `joint_distribution` computes it by enumerating every assignment of the variables and multiplying the mechanisms' entries, which makes it an oracle for small models rather than a scalable backend – `BayesianNetworkInference.jl` has the algorithms that scale. The result is a `FiniteKernel` state $I \rightarrow \text{Climate} \otimes \dots \otimes \text{Occupancy}$, so it carries its own axes.

```
J = joint_distribution(m)
J.codom
```

```
FiniteSpace(Climate{dry,normal,wet} ⊗ Irrigation{low,high} ⊗
SoilMoisture{low,medium,high} ⊗ GrazingPressure{low,high} ⊗
Vegetation{sparse,moderate,dense} ⊗ HabitatQuality{poor,good} ⊗
Occupancy{absent,present})
```

`joint_table` presents the same numbers as a table indexed by assignments:

```
jt = joint_table(m)
jt[:Climate ⇒ :dry, :Irrigation ⇒ :low, :SoilMoisture ⇒ :low,
   :GrazingPressure ⇒ :low, :Vegetation ⇒ :sparse, :HabitatQuality ⇒ :poor,
   :Occupancy ⇒ :absent]
```

0.025704

Marginals and conditionals are the usual queries, computed from that joint. Both take evidence (a model's own evidence by default), and `conditional` returns a kernel rather than a number, so a single entry is read off with probability; extracting such a kernel from a joint state is disintegration ([Cho and Jacobs 2019](#)).

```
round.(marginal(m, :Occupancy).table; digits = 4)
```

```
2-element Vector{Float64}:
 0.5238
 0.4762
```

```
probability(conditional(m, :Occupancy, :Vegetation), :present, :dense)
```

0.6675

sample draws ancestrally in topological order, each variable from its mechanism's kernel given the sampled parents. With enough draws the empirical frequencies approach the marginal, which is a cheap independent check on the enumeration:

```
using Random

draws = sample(m, 20_000; rng = MersenneTwister(2026))
emp = empirical_marginal(draws, :Occupancy, space(m, :Occupancy))
round.(emp.table; digits = 3), round.(marginal(m, :Occupancy).table; digits = 3)
```

```
([0.524, 0.476], [0.524, 0.476])
```

```
maximum(abs.(emp.table .- marginal(m, :Occupancy).table)) < 0.02
```

```
true
```

9 Summary

A network here is variables, states and mechanisms stored as an ACSet; the DAG, parents, children and topological_order are all views derived from the Input parts, and validate checks the structural rules in one place. Syntax carries no numbers: a BayesModel binds a FiniteKernel per mechanism, so the same structure can carry different tables or none, and joint_distribution, marginal, conditional and sample read the semantics off the bound model by brute force. The next vignette, [Observation versus intervention](#), separates conditioning from the do operator. The categorical view – the same structure as a string diagram, evaluated in FinStoch – is the first vignette of CategoricalBayesianNetworks.jl.

10 References

- Cho, Kenta, and Bart Jacobs. 2019. “Disintegration and Bayesian Inversion via String Diagrams.” *Mathematical Structures in Computer Science* 29 (7): 938–71. <https://doi.org/10.1017/S0960129518000488>.
- Fong, Brendan. 2012. *Causal Theories: A Categorical Perspective on Bayesian Networks*. <https://arxiv.org/abs/1301.6201>.
- Koller, Daphne, and Nir Friedman. 2009. *Probabilistic Graphical Models: Principles and Techniques*. MIT Press.
- Patterson, Evan, Owen Lynch, and James Fairbanks. 2022. “Categorical Data Structures for Technical Computing.” *Compositionality* 4 (5). <https://doi.org/10.32408/compositionality-4-5>.