

Dynamic networks: templates and unrolling

Simon Frost

Table of contents

1	Overview	1
2	Setup	2
3	The template	2
4	Unrolling	3
5	A lag-three boundary is three initial slices	4
6	Unrolling is gluing	5
7	Semantics	5
8	Forward marginals	6
9	Intervening in one slice	7
10	Management as an exogenous input	8
11	Summary	8
12	References	8

1 Overview

Ecological feedback is a cycle: vegetation feeds herbivores and herbivores eat vegetation. A Bayesian network cannot hold a cycle, but a *dynamic* Bayesian network can, by placing the two halves of the loop in different time slices: $\text{Vegetation}_t \rightarrow \text{Herbivores}_t$ within a slice and $\text{Herbivores}_{\{t-1\}} \rightarrow \text{Vegetation}_t$ across slices, which is the standard two-slice presentation of a dynamic Bayesian network (Koller and Friedman 2009). `BayesianNetworks.jl` represents such a model as a *template* (`DynamicBayesNet`) that `unroll` compiles, for any finite horizon, into an ordinary closed network. Every facility of the package then applies unchanged: validation,

evaluation, evidence, interventions, drawing and, through `BayesianNetworkInference.jl`, variable elimination on long horizons – the operation whose categorical form is marginalisation in a Markov category (Lorenzin and Zanasi 2025).

2 Setup

```
using BayesianNetworks
```

3 The template

A template has two parts, both plain BayesNets. The *transition* template describes a generic slice t : its *current* variables (Vegetation, Herbivores) have mechanisms, and its *lagged* variables, named with the convention `lagged(:X, k)` ($X[t-1]$, $X[t-2]$, ...), are exogenous stand-ins for earlier slices. The *initial* network is a closed network over the first slice. The example below is `vegetation_herbivore_dbn()`.

```
initial = bayesnet(:Vegetation => [:sparse, :dense], :Herbivores => [:low, :high];
               mechanisms = [:Herbivores => :Vegetation])
transition = bayesnet(:Vegetation => [:sparse, :dense], :Herbivores => [:low,
:high],
                    lagged(:Vegetation, 1) => [:sparse, :dense],
                    lagged(:Herbivores, 1) => [:low, :high];
                    mechanisms = [:Vegetation => (lagged(:Vegetation, 1),
                                                lagged(:Herbivores, 1)),
                                :Herbivores => :Vegetation],
                    closed = false)
dbn = DynamicBayesNet(initial, transition)
```

```
DynamicBayesNet(2 variables per slice, 2 lagged inputs, lags = 1)
```

```
current_variables(dbn), lagged_variables(dbn), lags(dbn)
```

```
([:Vegetation, :Herbivores], [Symbol("Vegetation[t-1]"), Symbol("Herbivores[t-1]")],
1)
```

The constructor validates the template: every current variable needs a mechanism, lagged variables must not have one, every $X[t-k]$ must be a lag of a current X with the same states, and the initial network must cover exactly the current variables.

```

bad = bayesnet(:Vegetation => [:sparse, :dense], :Herbivores => [:low, :high],
              lagged(:Vegetation, 1) => [:sparse, :dense];
              mechanisms = [:Vegetation => lagged(:Vegetation, 1)], closed = false)
try
  DynamicBayesNet(initial, bad)
catch e
  sprint(showerror, e)
end

```

"DynamicTemplateError (missing_mechanism): variable :Herbivores: current variable of the transition has no mechanism"

4 Unrolling

`unroll(dbn, h)` builds the closed network over slices 0, ..., h. Variables are named `variable_at(:X, t)`, that is `X_t`; slice 0 comes from the initial network and every later slice from the transition template, with each lagged input `X[t-k]` wired to `X_{t-k}`. The boundary is exact: with maximal lag `L` the initial network covers slices 0 to `L - 1` (for `L = 1`, just slice 0), so no lagged parent is ever missing.

```

u = unroll(dbn, 3)
validate(u; closed = true, unique_names = true)
variable_names(u)

```

```

8-element Vector{Symbol}:
 :Vegetation_0
 :Herbivores_0
 :Vegetation_1
 :Herbivores_1
 :Vegetation_2
 :Herbivores_2
 :Vegetation_3
 :Herbivores_3

```

```
variable_name.(Ref(u), parents(u, :Vegetation_2))
```


```

2-element Vector{Symbol}:
 :Vegetation_1
 :Herbivores_1

```

`Vegetation_2` reads slice 1, so the lag has become an ordinary arrow between slices, which the drawing lays out left to right:

```
to_graphviz(u; rankdir = "LR")
```



`slice_variables` and `horizon` read the slice structure back from the names:

```
slice_variables(u, 2), horizon(u)
```

```
([:Vegetation_2, :Herbivores_2], 3)
```

5 A lag-three boundary is three initial slices

For maximal lag three, the initial network supplies slices 0, 1 and 2 jointly. Its lag notation is relative to slice 2: $X[t-2]$, $X[t-1]$ and X become X_{-0} , X_{-1} and X_{-2} . They need not have the same initial distribution.

```
initial3 = bayesnet(lagged(:X, 2) => [:no, :yes],
                    lagged(:X, 1) => [:no, :yes], :X => [:no, :yes])
transition3 = bayesnet(:X => [:no, :yes], lagged(:X, 3) => [:no, :yes];
                    mechanisms = [:X => lagged(:X, 3)], closed = false)
dm3 = DynamicBayesModel(DynamicBayesNet(initial3, transition3))
dm3 = bind_cpt(dm3, [lagged(:X, 2) => [0.9, 0.1],
                    lagged(:X, 1) => [0.5, 0.5], :X => [0.2, 0.8]]; slice =
                    :initial)
dm3 = bind_cpt(dm3, :X => [0.8 0.2; 0.1 0.9])
u3 = unroll(dm3, 5)
[(slice = t,
  parents = variable_name.(Ref(syntax(u3)), parents(syntax(u3), variable_at(:X,
t))),
  probability_yes = round(marginal(u3, variable_at(:X, t)).table[2]; digits = 4))
 for t in 0:5]
```

```
6-element Vector{NamedTuple{(:slice, :parents, :probability_yes)}}:
 (slice = 0, parents = Any[], probability_yes = 0.1)
 (slice = 1, parents = Any[], probability_yes = 0.5)
 (slice = 2, parents = Any[], probability_yes = 0.8)
 (slice = 3, parents = [:X_0], probability_yes = 0.27)
 (slice = 4, parents = [:X_1], probability_yes = 0.55)
 (slice = 5, parents = [:X_2], probability_yes = 0.76)
```

The first transition is at slice 3, reading slice 0, not an invented negative slice or a repeated copy of the slice-2 initial state. Since the transition sends a probability p to $0.2 + 0.7p$, slices 3, 4 and 5 have probabilities 0.27, 0.55 and 0.76. Dropping the missing lag or reusing one initial marginal for all three slices would change those numbers.

6 Unrolling is gluing

The unrolled network is what the open-network machinery produces when consecutive slices are glued along the lagged variables, as structured cospans (Baez and Courser 2020). `initial_slice` is the initial network renamed to slice 0 and `transition_slice(dbn, t)` the transition template seen from slice `t`, with the lagged variables exogenous.

```
variable_names(initial_slice(dbn)), variable_names(transition_slice(dbn, 1))
```

```
([:Vegetation_0, :Herbivores_0], [:Vegetation_1, :Herbivores_1, :Vegetation_0, :Herbivores_0])
```

Gluing the outputs of the first to the inputs of the second gives `unroll(dbn, 1)` up to renumbering. `Open` and `glue` live in `CategoricalBayesianNetworks.jl`, which shows that identity as a worked example; `unroll` builds the network directly rather than by repeated pushouts, and both test suites keep the two constructions in agreement.

7 Semantics

A `DynamicBayesModel` carries a `BayesModel` for each template. Kernels are bound with `bind_cpt` (or `bind_kernel`) on the chosen slice; the domain axes of a transition kernel are named like the template's inputs, lagged names included. `vegetation_herbivore_model()` binds fixed tables: vegetation recovers when herbivores are few and declines when they are many, and herbivores track vegetation.

```
dm = vegetation_herbivore_model()
kernel(dm, :Vegetation)
```

```
FiniteKernel{Float64}(Vegetation[t-1]{sparse,dense} ⊗ Herbivores[t-1]{low,high} →
Vegetation{sparse,dense})
      sparse,low  dense,low  sparse,high  dense,high
sparse      0.5      0.1      0.8      0.4
dense      0.5      0.9      0.2      0.6
```

```
round.(kernel(dm, :Vegetation).table; digits = 2) # outputs-first: (Vegetation,
Vegetation[t-1], Herbivores[t-1])
```

```
2×2×2 Array{Float64, 3}:
```

```
[:, :, 1] =
 0.5  0.1
 0.5  0.9
```

```
[:, :, 2] =
 0.8  0.4
 0.2  0.6
```

unroll on the model copies the kernels into every slice, renaming their axes, so the result is an ordinary model with semantics.

```
um = unroll(dm, 3)
validate(um; closed = true, unique_names = true, semantics = true)
names(kernel(um, :Vegetation_2).dom)
```

```
2-element Vector{Symbol}:
 :Vegetation_1
 :Herbivores_1
```

8 Forward marginals

rollout returns the marginal of a variable in every slice (brute force through marginal, fine for short horizons); filter_marginal asks for one slice, given the model's evidence.

```
rv = rollout(um, :Vegetation)
rh = rollout(um, :Herbivores)
fmt(x) = lpad(string(round(x; digits = 4)), 8)
println(" t    P(Vegetation = dense)    P(Herbivores = high)")
for t in 0:horizon(syntax(um))
    println(lpad(t, 2), " ", lpad(fmt(rv[t + 1].table[2]), 21), " ",
            lpad(fmt(rh[t + 1].table[2]), 20))
end
```

t	P(Vegetation = dense)	P(Herbivores = high)
0	0.6	0.54
1	0.578	0.5312
2	0.5718	0.5287
3	0.5701	0.528

The same numbers follow from the forward recursion of the two-variable chain: the joint of (Vegetation_t, Herbivores_t) composed with the transition kernel gives the next vegetation marginal. Kernels compose, so this is one line.

```
J0 = marginal(um, [:Vegetation_0, :Herbivores_0])
T = kernel(um, :Vegetation_1)
compose_kernel(J0, T).table ≈ rv[2].table
```

```
true
```

Evidence in one slice propagates both forwards (filtering) and backwards (smoothing):

```
ume = observe(um, :Herbivores_1 ⇒ :high)
(round.(filter_marginal(ume, :Vegetation, 2).table; digits = 4),
 round.(filter_marginal(ume, :Vegetation, 0).table; digits = 4))
```

```
([0.4953, 0.5047], [0.3494, 0.6506])
```

For long horizons the unrolled model is just a large network: `BayesianNetworkInference.infer` answers the same queries by variable elimination without enumerating the joint.

9 Intervening in one slice

An intervention on `Herbivores_2` is a local rewrite of the unrolled network: the mechanism of that one variable is replaced by a point mass. Earlier slices are untouched and later ones respond.

```
umi = do_intervention(um, :Herbivores_2 => :high)
rvi = rollout(umi, :Vegetation)
println(" t    P(dense)    P(dense | do(Herbivores_2 = high))")
for t in 0:3
    println(lpad(t, 2), "    ", fmt(rv[t + 1].table[2]), "    ", fmt(rvi[t + 1].table[2]))
end
```

t	P(dense)	P(dense do(Herbivores_2 = high))
0	0.6	0.6
1	0.578	0.578
2	0.5718	0.5718
3	0.5701	0.4287

```
to_graphviz(umi; rankdir = "LR", states = false)
```



Observation of the same event, by contrast, also moves the earlier slices:

```
rvo = rollout(observe(um, :Herbivores_2 => :high), :Vegetation)
round.([k.table[2] for k in rvo]; digits = 4)
```

```
4-element Vector{Float64}:
 0.6142
 0.6297
 0.7571
 0.5028
```

10 Management as an exogenous input

`vegetation_herbivore_dbn(management = true)` adds an exogenous Management variable (none, cull) to every slice, feeding the herbivore mechanism. Setting it in one slice by evidence or intervention gives the same downstream effect, because it has no parents.

```
umm = unroll(vegetation_herbivore_model(management = true), 3)
(round.(marginal(umm, :Vegetation_3).table; digits = 4),
 round.(marginal(umm, :Vegetation_3; evidence = Dict(:Management_2 => :cull)).table;
 digits = 4),
 round.(marginal(do_intervention(umm, :Management_2 => :cull), :Vegetation_3).table;
 digits = 4))
```

```
([0.4088, 0.5912], [0.3466, 0.6534], [0.3466, 0.6534])
```

11 Summary

- A DynamicBayesNet is an initial network plus a transition template whose lagged inputs follow the $X[t-k]$ naming convention; `validate` checks the template.
- `unroll(dbn, h)` compiles a finite horizon into a closed BayesNet with variables X_t ; on a DynamicBayesModel the kernels are copied per slice.
- Unrolling agrees with gluing consecutive slices as open networks.
- `rollout` and `filter_marginal` give forward marginals by brute force; evidence and interventions on individual slices work as in any other model.

The next vignette, [Model cards and provenance](#), documents a finished model rather than computing with it.

12 References

- Baez, John C., and Kenny Courser. 2020. “Structured Cospans.” *Theory and Applications of Categories* 35 (48): 1771–822.
- Koller, Daphne, and Nir Friedman. 2009. *Probabilistic Graphical Models: Principles and Techniques*. MIT Press.
- Lorenzin, Antonio, and Fabio Zanasi. 2025. *Bayesian Networks, Markov Networks, Moralisation, Triangulation: A Categorical Perspective*. <https://arxiv.org/abs/2512.09908>.